

# Chapter 10

DJM

10 March 2020

## Quick review of OLS

- OLS stands for “ordinary least-squares”.
- Essentially, it means “solve the least-squares problem”

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^n (x_i^\top \beta - y_i)^2 = (X^\top X)^{-1} X^\top Y$$

- The hat matrix is

$$\hat{Y} = X\hat{\beta} = X(X^\top X)^{-1} X^\top Y = HY$$

- The Gauss-Markov theorem says if:
  1.  $Y_i = x_i^\top \beta + \epsilon_i$
  2.  $\mathbb{E}[\epsilon_i] = 0$
  3.  $\mathbb{V}[\epsilon_i] = \sigma^2 < \infty$
  4.  $\operatorname{Cov}[\epsilon_i, \epsilon_j] = 0$

Then  $\hat{\beta}$  has the smallest variance of all possible unbiased estimators for  $\beta$ .

## What is WLS and why use it?

- Weighted least-squares (WLS) is simply

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^n w_i (x_i^\top \beta - y_i)^2 = (X^\top W X)^{-1} X^\top W Y$$

- If some of those assumptions for G-M are violated, in particular, if  $\mathbb{V}[\epsilon_i]$  depends on  $x_i$  (notated like  $\sigma^2(x_i)$ ), then we lose the optimality.
- Aside: Gauss-Markov is a commonly used justification for OLS in applied work. The logic goes like this: (1) unbiased is good, (2) G-M says OLS is the best linear model which is unbiased. The problem is that (1) is wrong. Unbiased may be good, but often a little bias is better.
- So what does WLS do?
  1. It **is** optimal, in the sense of G-M, if  $\mathbb{V}[\epsilon_i] = \sigma_i^2$ .
  2. You’ve already used it (see next slide).
  3. What if you want to predict  $Y_i$  which have other structures like  $y_i \in \{0, 1\}$ ? The algorithms for the new estimators (often called GLS for generalized least squares) use WLS (logistic regression is one example we are building to).

## You already used WLS

- I said you already did this. It is convenient that Kernel regression **is** WLS.
- In particular Kernel regression looks like

$$\hat{c} = \underset{c}{\operatorname{argmin}} \sum_{j=1}^n \sum_{i=1}^n w_{ij} (c_j - y_i)^2 \quad w_{ij} = \frac{K((x_i - x_j)/h)}{\sum_{i=1}^n K((x_i - x_j)/h)}$$

This is locally constant regression.

- You don't need to understand this formula, but it can be useful, and it provides some justification for WLS based on previous ideas.

## What goes wrong with heteroskedasticity?

- So suppose  $\mathbb{V}[\epsilon_i] = \sigma^2(x_i)$ . That is our “homoskedasticity” assumption is violated. Should we care?
- What if we just use OLS (that is `lm`) anyway?
- Some things don't change.
  1. We still have that  $\mathbb{E}[\hat{\beta}] = \beta$ . That is OLS **is** still unbiased.
  2. We still have that OLS minimizes the sum of squared residuals: among all lines, OLS makes  $\sum_{i=1}^n (x_i^\top \hat{\beta} - y_i)^2$  as small as possible.
- Some things **do** change.
  1. OLS no longer has the best variance of all unbiased estimators (WLS does).
  2. The standard errors that R produces are wrong. They make it seem “more certain” than is correct (could use the bootstrap to fix it though).
  3. So are the  $F$ -tests and  $p$ -values (again, the bootstrap).

## Log squared residuals

- So WLS is fairly general. But for now, let's focus on how to use it for heteroskedasticity.
- Suppose you **know** the following:
  1.  $Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$ .
  2. You know  $\beta_0 = 3$  and  $\beta_1 = 2$ .
  3.  $\mathbb{E}[\epsilon_i] = 0$
  4.  $\mathbb{V}[\epsilon_i] = \sigma^2(x_i)$  ( $\sigma^2(\cdot)$  is a function).
- You don't know  $\sigma^2(\cdot)$
- How would you estimate  $\sigma^2(x)$ ?
- You can use nonparametric regression of course!
  - Just look at  $e_i = y_i - 3 + 2x_i$ .
  - We already know that  $e_i$  has mean zero, so no use estimating it's mean
  - Therefore  $\mathbb{E}[e_i^2] = \sigma^2(x_i)$ .
  - Now this is easy: rewrite the model as  $e_i^2 = \sigma^2(x_i) + \eta_i$
  - It's just nonparametric regression (since you know  $e_i^2$  and  $x_i$ ).

## So why not?

- There's one problem with the above line of reasoning:  $e_i^2 > 0$  so there are some constraints on  $\eta_i$ .
- Imagine if, say, at  $x = 1$ ,  $\sigma^2(1) = .001$ , then it's pretty likely that  $\eta$  is large and positive there, so there's heteroskedasticity in our model for heteroskedasticity...

- If the original  $\epsilon_i$  were  $N(0, \sigma^2(x_i))$ , then the new  $\eta_i$  are distributed as  $\chi_1^2$ , so these are right skewed, everywhere.
- A remedy is to look at  $\log e_i^2 = \log \sigma^2(x_i) + \tau_i$ .
- You could try both and look at qq-plots of the residuals. I tend to prefer the second set-up. It's just more satisfying.
- This is just a transformation: just like when you looked at qq-plots and decided to model  $\log Y$  rather than  $Y$ .

## What's the Oracle?

- So back to WLS.
- I had that example where I wanted to estimate the variance function
- In that case **I knew the mean:  $3 + 2x$**
- I knew because the Oracle told me. The Oracle is a wise woman who lives at Delphi according to the ancient Greeks, and speaks the thoughts of Apollo.
- In other words, she tells me things no one could possibly know, like the mean function.
- In statistics, we talk about Oracles a lot, usually as a way of comparing a procedure we cook up for estimating something to the answers the Oracle would have told us (the best possible, but unobtainable estimator).

## A big example

- This is a (slightly modified) portion of a real job interview.
- It is a very simple application of heteroskedasticity.
- Heteroskedasticity appears frequently with financial data, so those companies like to see if you can handle it.

## The set up

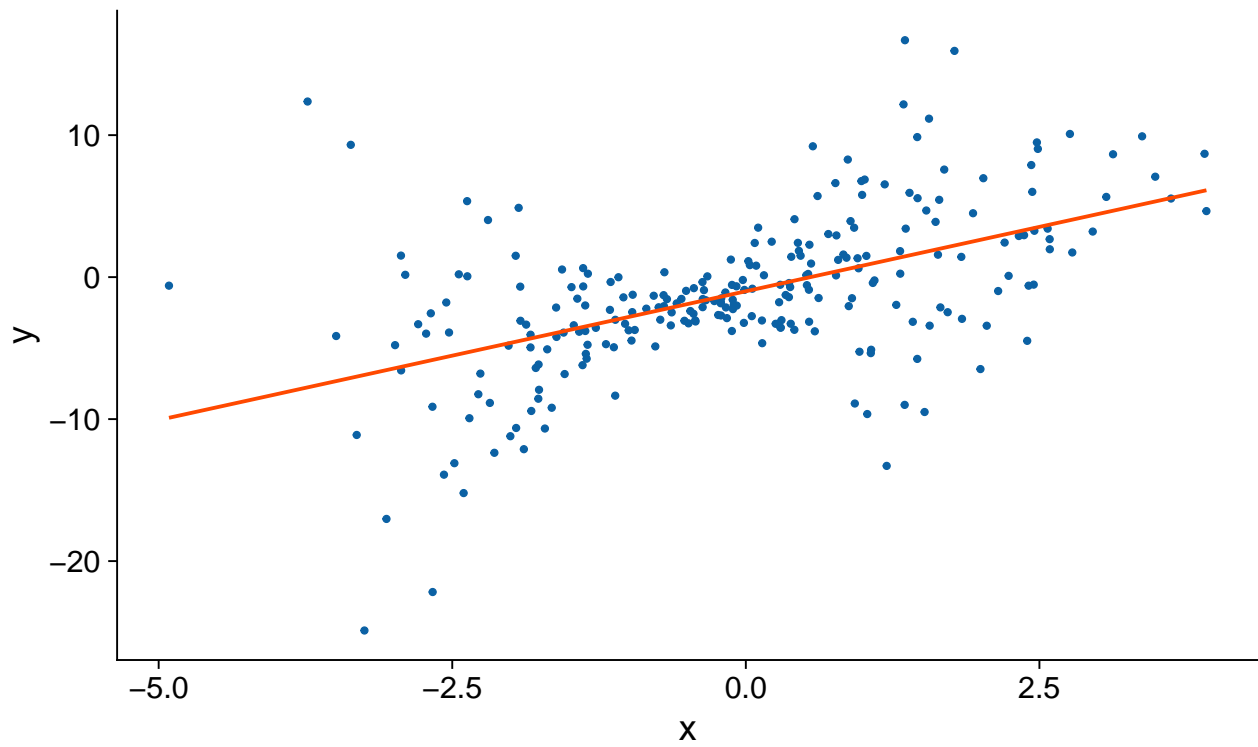
- The dataset `jobInt` contains data from a simple linear model with heteroskedastic noise.
- In other words, for  $i = 1, \dots, 250$ ,

$$y_i = \beta_0 + \beta_1 x_i + \sigma(x_i) \epsilon_i \quad \epsilon_i \sim N(0, 1).$$

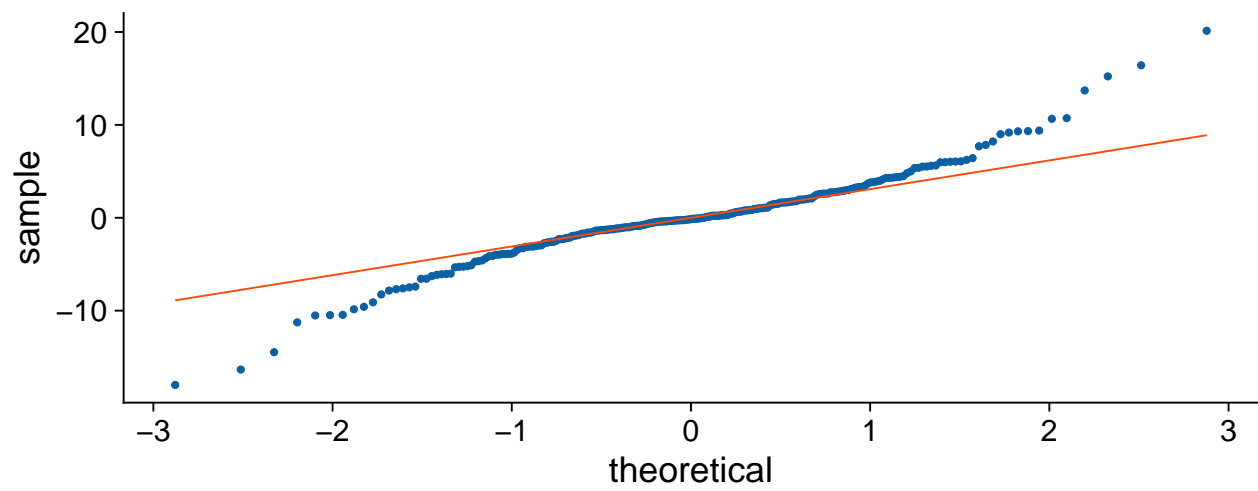
- You know nothing about (the function)  $\sigma(\cdot)$ .
- Your goal is to estimate  $(\beta_0, \beta_1)$  as well as possible, and provide a CI.

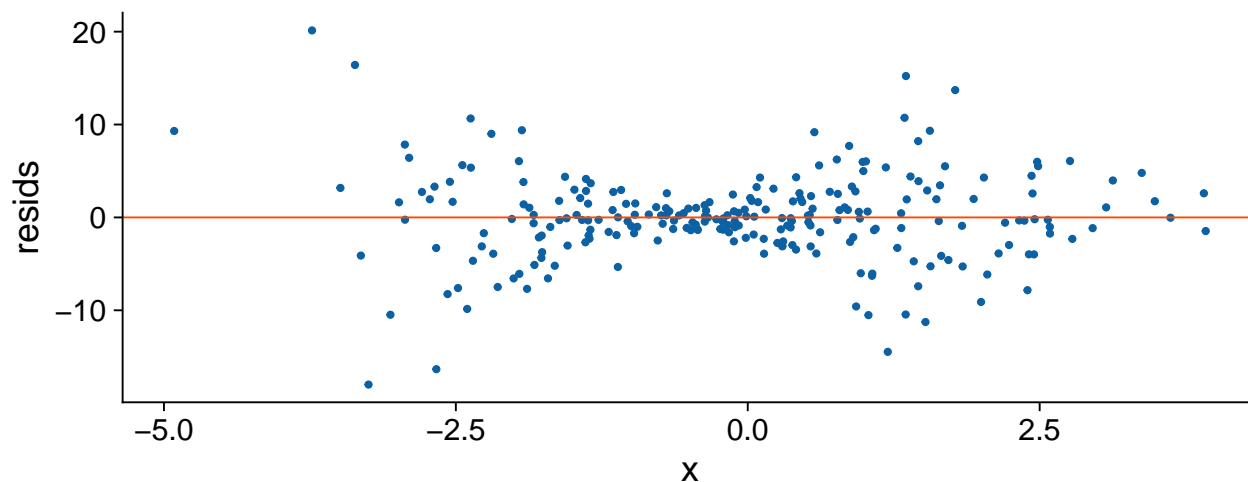
## How do I do this?

- First things first, EDA.

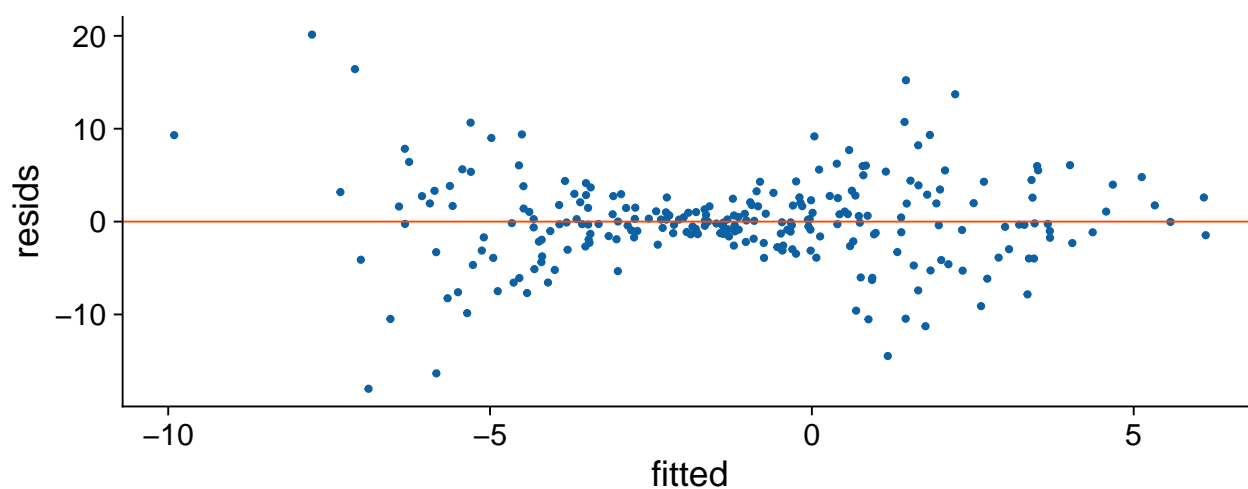


qq-plot and residuals against x





residuals vs. fitted values



So now what?

- We **know** that  $\mathbb{E}[Y \mid X = x] = \beta_0 + \beta_1 x$ .
- We **know** that  $\mathbb{E}[\hat{e} \mid X = x] = 0$ .
- We **know** that  $\mathbb{E}[\hat{e}^2 \mid X = x] = \sigma^2(x)$ .
- So we want to try to estimate  $\sigma^2(x)$  and  $\beta_0$  and  $\beta_1$  all at the same time.

Oracle information

- If we knew  $\beta_0$  and  $\beta_1$ , then we could use `npreg` to estimate  $\sigma^2(x)$ .
- If we **knew**  $\sigma^2(x)$ , then we could use WLS to estimate  $\beta_0$  and  $\beta_1$  (and all the SEs would be right!)
- But we don't know either.

Procedure

1. Use `lm` to estimate  $\beta_0$  and  $\beta_1$ .
2. Now pretend that you “know” them, calculate  $\log(\hat{e}^2)$  and use `npreg` to estimate  $\log \sigma^2(x)$ .

- Now pretend that you “know”  $\sigma^2(x)$  (take exp of your estimate from 2.) and use WLS (with `lm(y~x, weights=1/sig2)`)
- You could stop here. But since you now have “better” estimates of  $\beta_1$  and  $\beta_0$ , it’s better to iterate 2 and 3 until some convergence.
- Ok. Something converged, so you return the last estimates of  $\beta_0$  and  $\beta_1$ . But the SEs are not right (because you “know”  $\sigma^2(x)$  but you don’t **know** it).
- To get SEs, use the bootstrap:
  - Non-parametric: repeat 1-5  $B$  times on resampled data.
  - Model-based: this is actually pretty hard here, better not to do it.

## Some code

- This code takes in data and does steps 1-5. It is **not** optimized for speed, but for readability, so run with care.


```
heteroWLS <- function(dataFrame, tol = 1e-4, maxit = 100, track=FALSE){
  # inputs: a data object, optional: tolerance, max.iterations, and progress tracker (prints)
  # outputs: estimated betas and weights
  require(np)
  ols = lm(y~x, data=dataFrame)
  b = coefficients(ols)
  conv = FALSE
  for(iter in 1:maxit){ # don't let this run forever
    if(conv) break # if the b's stop moving, get out of the loop
    logSqResids = log(residuals(ols)^2)
    winv = exp(predict(npreg(logSqResids~x, data=dataFrame, tol=1e-2, ftol=1e-2)))
    winv[winv < tol] = tol # zero inverse weights are bad, make them small
    ols = lm(y~x, weights = 1/winv, data=dataFrame) #weights are 1 / estim.variance
    newb = coefficients(ols)
    conv.crit = sum((b-newb)^2) # calculate how much b moved
    if(track) cat('\n', iter, '/', maxit, ' conv.crit = ', conv.crit) # print progress
    conv = (conv.crit < tol) # check if the b's changed much
    b = newb # update the coefficient estimates
  }
  return(list(betas=b, weights = winv, log2resids = log(residuals(ols)^2)))
}
```

## Do it! (takes a little while...)

```
resampWLS <- function(dataFrame,...){ # ... means options passed on
  rowSamp = sample(1:nrow(dataFrame), size=nrow(dataFrame), replace=TRUE)
  return(heteroWLS(dataFrame[rowSamp,],...)$betas) # passed things on if desired
}
B = 100 #
alp = .05
origBetas = heteroWLS(jobInt)
system.time(bootBetas <- replicate(B, resampWLS(jobInt, maxit=20)))
qq = apply(bootBetas, 1, quantile, probs=c(1-alp/2, alp/2))
CI = cbind(origBetas$betas, 2*origBetas$betas - t(qq))
colnames(CI) = c('coef', rev(colnames(CI)[2:3]))
```

```
##              coef      2.5%      97.5%
## (Intercept) -0.9938074 -1.406956 -0.4306921
## x           2.0023640  1.607402  2.6352379
```

## Some plots



## Some caveats

- If we **care** about estimating  $\sigma(\cdot)$ , then what we did is ok.
- But we don't.
- We only care about estimating  $\beta_0$  and  $\beta_1$ .
- So better to use CV with leaving out  $(y_i, x_i)$ , instead of using CV to estimate  $\sigma(\cdot)$  (which is what `npreg` is doing; it knows nothing about  $y_i$ )
- This takes a bit more work to code up (Try it!)

## Local linear vs. Kernels


- People often wonder whether to use local linear regression or Kernels.
- Like with many things, there isn't really a cut-and-dried answer.
- Some practitioners prefer local linear regression.
- It's main benefit is to correct "boundary bias".
- Otherwise, not much different.

## Repeat of Ch. 4

- We can estimate this easily with both a kernel and local linear regression

```
kern = npreg(y~x, data = df)
locclin = npreg(y~x, data=df, regtype='ll')
df$kern = fitted(kern)
df$locclin = fitted(locclin)
df$truth = trueFunction(x)
dflines = pivot_longer(select(df,-y),-x, names_to = 'method')
ggplot(df, aes(x, y)) + geom_point() + xlim(0,2*pi) + ylim(0,max(df$y)) +
```

```
geom_line(data=dflines, aes(x=x,y=value,color=method)) +  
scale_color_manual(values=c(green,blue,red))
```



## What is Loess?

- So kernel regressions are local constants, we then saw local linear regression which is quite similar.
- Why stop there? The next term is squared things, then cubics, then...
- Loess uses local polynomials (of some order) in a particular way (combining  $k$ -nearest neighbor regression with subsampling).
- It is actually quite cool, but its complexity makes it hard to deal with.
- Theoretically, one can show that Kernels are optimal, so it's not really worth worrying too much about, but it can work well, and it doesn't require installing a package.
- Here it is on my previous example

```
ls = loess(y~x, data=df)  
df$loess = fitted(ls)  
dflines = pivot_longer(select(df,-y),-x,names_to='method')  
ggplot(df, aes(x, y)) + geom_point() + xlim(0,2*pi) + ylim(0,max(df$y)) +  
  geom_line(data=dflines, aes(x=x,y=value,color=method)) +  
  scale_color_manual(values=c(green,blue,red,orange))
```



